

In [2]:

```
"""  
Arjun Srivastava  
arj1  
AMATH 301 B  
"""  
  
import numpy as np  
import scipy.linalg  
import pandas as pd  
import time
```

```
In [3]: # Problem 1

# a)

a, b = 1e21 - (1e21 - 1e5), (1e21 - 1e21) + 1e5

diff = abs(a-b) # diff = 31072.0

"""
Considering the size of the numbers a and b, I would initially say that this r
esult is not very significant. 31072.0, however,
is still a large number, and there are many contexts in which this error would
be significant.
"""

# b)

A = np.array([[10e-20, 1], [1, 1]])

cond = np.linalg.cond(A) # cond = 2.6180339887498953, which implies a unique
solution

# c)

L, U = np.array([[1, 0], [10e20, 1]]), np.array([[10e-20, 1], [0, 1-10e20]])
print('='*30, 'c)', sep='\n')
print(A, (L @ U), sep='\n\n')

"""
L @ U is not equal to A. The first row is, and this is because we are not addi
ng values that involve catastrophic cancellation.
The bottom row, however, involves 10e20 + 10e-20, which should be equal to 0.
As can be seen by the result printed below,
there is an extra 100. This makes sense due to the immense size of 10e20 and 1
0e-20. The second column of the second row
has 10e20 - 10e20. The result should be 1, but it is instead 0. Catastrophic c
ancellation is occurring in both of these cases
since 10e20 and 10e-20 are being operated on together.
"""

# d)

B = np.array([[1, 1], [10e-20, 1]])
L1, U1 = np.array([[1, 0], [10e-20, 1]]), np.array([[1, 1], [0, 1-10e-20]])
print('='*30, 'd)', sep='\n')
print(B, (L1 @ U1), sep='\n\n')
print('='*30, 'e)', sep='\n')

"""
By reordering the matrix A to B, we have avoided any operation that involves t
wo very large numbers (10e20, 10e-20). Since
10e20 and 10e-20 are only ever being added to or subtracted by 1, there is no
case in which catastrophic cancellation occurs,
and our answers are much more accurate.
"""
```

```
# e)

P, L2, U2 = scipy.linalg.lu(A)
P = P.T
print(L2, U2, P, sep='\n')
print('\n')

P1, L3, U3 = scipy.linalg.lu(B)
P1 = P1.T
print(L3, U3, P1, sep='\n')

"""
The official SciPy LU yields the same results as part d). This implies that Python is using the permutation matrix to ensure that the matrix A can be more efficiently decomposed without issues. As shown, P is different for A and B, implying that Python wants to avoid catastrophic cancellation by setting up the equation in a more efficient manner. Since A is not set up efficiently for an LU decomp, the permutation matrix converts it.
"""
```

```

=====
c)
[[1.e-19 1.e+00]
 [1.e+00 1.e+00]]

[[1.e-19 1.e+00]
 [1.e+02 0.e+00]]
=====
d)
[[1.e+00 1.e+00]
 [1.e-19 1.e+00]]

[[1.e+00 1.e+00]
 [1.e-19 1.e+00]]
=====
e)
[[1.e+00 0.e+00]
 [1.e-19 1.e+00]]
[[1. 1.]
 [0. 1.]]
[[0. 1.]
 [1. 0.]]

[[1.e+00 0.e+00]
 [1.e-19 1.e+00]]
[[1. 1.]
 [0. 1.]]
[[1. 0.]
 [0. 1.]]

```

Out[3]: '\n\nThe official SciPy LU yields the same results as part d). This implies that Python is using the permutation matrix\nto ensure that the matrix A can be more efficiently decomposed without issues. As shown, P is different for A and B,\nimplying that Python wants to avoid catastrophic cancellation by setting up the equation in a more efficient manner. Since\nA is not set up efficiently for an LU decomp, the permutation matrix converts it.\n'

In [4]: *# Problem 2*

a)

```
A = np.genfromtxt('example_matrix.csv', delimiter=',')
```

In [5]: *# b)*

```

t0 = time.time()
rb = []
for i in range(100):
    b = np.random.rand(3000, 1)
    x = scipy.linalg.solve(A, b)
    rb.append(np.max(np.abs((A @ x) - b)))
t1 = time.time()
final_time_b = t1 - t0 # 44.55109429359436 seconds
rb = np.mean(rb) # 2.954359e-14

```

```
In [6]: # c)

t0 = time.time()
P, L, U = scipy.linalg.lu(A)
P = P.T
rc = []
for i in range(100):
    b = np.random.rand(3000, 1)
    y = scipy.linalg.solve_triangular(L, P @ b, lower=True)
    x = scipy.linalg.solve_triangular(U, y)
    rc.append(np.max(np.abs((A @ x) - b)))
t1 = time.time()
final_time_c = t1 - t0 # 4.096189737319946 seconds
rc = np.mean(rc) # 2.958356e-14
```

```
In [11]: # d)

t0 = time.time()
inv = scipy.linalg.inv(A)
rd = []
for i in range(100):
    b = np.random.rand(3000, 1)
    x = inv @ b
    rd.append(np.max(np.abs((A @ x) - b)))
t1 = time.time()
final_time_d = t1 - t0 # 1.173656940460205 seconds
rd = np.mean(rd) # 9.092505e-14

# e) modified the code above ^
```

```
In [8]: # f)

data = {'Method': ['G.E.', 'LU', 'Inverse'], 'total time': [final_time_b, final_time_c, final_time_d], 'average residual': [rb, rc, rd]}
table = pd.DataFrame(data)

# g)

"""
If I only cared about speed, I would choose the inverse method, as it is significantly faster than the Gaussian Elimination method and fairly quicker than the LU decomposition. The LU decomposition is close behind (3 seconds slower).
"""

# h)

"""
If I only cared about accuracy, I would choose either the Gaussian Elimination method or the LU decomposition. They have near identical average residuals, and both are significantly lower than the inverse method.
"""

# i)

"""
The most ideal method seems like the LU decomposition. It has a fast time and much less residual error than the inverse method. Overall, it seems like the best tradeoff for guaranteeing fast and accurate results, as the Gaussian Elimination method takes significantly longer to execute.
"""

# j)

"""
The times are increased if these calculations are done inside the loop. For example, the inverse loop took 1.17 seconds before, but now takes 65.18 seconds. This is a very important detail, especially if we are trying to analyze the runtime of a specific function or process. The LU decomposition and the inverse calculations all take time, and we should always put pre and post steps outside of the time calculation. In the inverse case, the time increase was extremely significant, proving that all steps matter.
"""

# k)

"""
To better time the efficiency of the solve function, I would start and end my times() within the for loop itself, starting before the solve function and ending right after. After the time is closed for that particular iteration of the loop, I would add the difference between the ending time and starting time to an int
"""
```

(which was defined before the loop), and then check the value of the int after the for loop is done. This way, I am only spending time calculating the solve function, and not the creation of the range() object or the creation of the random vector. If I were to implement this, I am sure the times would be smaller across the board for each method.

"""

table # show DataFrame

Out[8]:

	Method	total time	average residual
0	G.E.	42.274952	2.954359e-14
1	LU	4.819105	2.958356e-14
2	Inverse	1.956770	9.092505e-14

In []: